

# STAT4060J hw4

Siqi Xiao

2022-11-01

## Question 1

1. Write the R code linear SVM model to classify the digits. I attached the data file and a file of code to be completed. Please monitor the training accuracy and testing accuracy and describe your finding.

```
library("data.table")
load_digits <- function(subset=NULL, normalize=TRUE) {
  #Load digits and labels from digits.csv.
  df <- fread("digits.csv")
  df <- as.matrix(df)
  ## only keep the numbers we want.
  if (length(subset)>0) {
    c <- dim(df)[2]
    l_col <- df[,c]
    index = NULL
    for (i in 1:length(subset)){
      number = subset[i]
      index = c(index,which(l_col == number))}
    sort(index)
    df = df[index,]
  }
  # convert to arrays.
  digits = df[,-1]
  labels = df[,c]
  # Normalize digit values to 0 and 1.
  if (normalize == TRUE) {
    digits = digits - min(digits)
    digits = digits/max(digits)}
  # Change the labels to 0 and 1.
  for (i in 1:length(subset)) {
    labels[labels == subset[i]] = i-1
  }
  return(list(digits, labels))}

split_samples <- function(digits,labels) {
  # Split the data into a training set (70%) and a testing set (30%).
  num_samples <- dim(digits)[1]
  num_training <- round(num_samples*0.7)
  indices = sample(1:num_samples, size = num_samples)
  training_idx <- indices[1:num_training]
  testing_idx <- indices[-(1:num_training)]
  return (list(digits[training_idx,], labels[training_idx,],
               digits[testing_idx,], labels[testing_idx]))}
```

```

result <- load_digits(subset=c(1, 7), normalize=TRUE)
digits = result[[1]]
labels = result[[2]]
result <- split_samples(digits,labels)
training_digits = result[[1]]
training_labels = result[[2]]
testing_digits = result[[3]]
testing_labels = result[[4]]
# print dimensions
dim(training_digits)

## [1] 253 64

dim(testing_digits)

## [1] 108 64

# Train a model and display training accuracy.
my_SVM <- function(X_train, Y_train, X_test, Y_test, lambda = 0.01, num_iterations = 1000, learning_rate = 0.01)
{
  n <- dim(X_train)[1]
  p <- dim(X_train)[2]+1
  X_train1 <- cbind(rep(1,n), X_train)
  Y_train <- 2*Y_train - 1#0,1->-1,1
  beta <- matrix(rep(0,p), nrow = p)

  ntest <- nrow(X_test)
  X_test1 <- cbind(rep(1,ntest), X_test)
  Y_test <- 2*Y_test - 1#0,1->-1,1

  acc_train <- rep(0, num_iterations)
  acc_test <- rep(0, num_iterations)


  for(it in 1:num_iterations){
    s <- X_train1 %*% beta
    db <- s*Y_train < 1
    dbeta <- matrix(rep(1,n), nrow = 1)%*%((matrix(db*Y_train, n, p)*X_train1))/n
    beta <- beta +learning_rate *t(dbeta)
    beta[2:p] <- beta[2:p] - lambda *beta[2:p]

    acc_train[it] <- mean(sign(s * Y_train))
    acc_test[it] <- mean(sign(X_test1 %*% beta * Y_test))
  }
  model <- list(beta = beta, acc_train = acc_train, acc_test = acc_test)
  model
}


SVM_model <- my_SVM(training_digits, training_labels, testing_digits, testing_labels, num_iterations = 1000)
beta <- SVM_model[[1]]
acc_train <- SVM_model[[2]]
acc_test <- SVM_model[[3]]

iter <- 1:200
plot(iter, acc_train, type = 'b',ylim = c(0,1),col = 'red')

```



```
plot(iter, acc_test, type = 'b', ylim = c(0,1), col = 'green')
```



From the result above, I find that the accuracy for both training and testing are climbing in a fast speed, and finally come to a perfect accuracy. The first iteration is the iteration the model improves the fastest.

## Question 2


2. Next, let's dive deeper into SVM, you should explore at least the following:
  - a. difference choice of  $C$  in controlling slackness
  - b. Kernel SVM with Gaussian radial basis function.
  - c. Another variation proposed by yourself.

```

library(plotrix)
model1 <- my_SVM(training_digits, training_labels, testing_digits, testing_labels, num_iterations = 10,
model2 <- my_SVM(training_digits, training_labels, testing_digits, testing_labels, num_iterations = 10,
model3 <- my_SVM(training_digits, training_labels, testing_digits, testing_labels, num_iterations = 10,
model4 <- my_SVM(training_digits, training_labels, testing_digits, testing_labels, num_iterations = 10,
model5 <- my_SVM(training_digits, training_labels, testing_digits, testing_labels, num_iterations = 10,
model6 <- my_SVM(training_digits, training_labels, testing_digits, testing_labels, num_iterations = 10,
model7 <- my_SVM(training_digits, training_labels, testing_digits, testing_labels, num_iterations = 10,
model8 <- my_SVM(training_digits, training_labels, testing_digits, testing_labels, num_iterations = 10,
model9 <- my_SVM(training_digits, training_labels, testing_digits, testing_labels, num_iterations = 10,
model10 <- my_SVM(training_digits, training_labels, testing_digits, testing_labels, num_iterations = 10,

iter <- 1:10
acc_train1 <- model1[[2]]
acc_test1 <- model1[[3]]
acc_train2 <- model2[[2]]
acc_test2 <- model2[[3]]
acc_train3 <- model3[[2]]
acc_test3 <- model3[[3]]
acc_train4 <- model4[[2]]
acc_test4 <- model4[[3]]
acc_train5 <- model5[[2]]
acc_test5 <- model5[[3]]
acc_train6 <- model6[[2]]
acc_test6 <- model6[[3]]
acc_train7 <- model7[[2]]
acc_test7 <- model7[[3]]
acc_train8 <- model8[[2]]
acc_test8 <- model8[[3]]
acc_train9 <- model9[[2]]
acc_test9 <- model9[[3]]
acc_train10 <- model10[[2]]
acc_test10 <- model10[[3]]
plot(iter,acc_train1,type = 'o',ylim = c(0,1),col = "red",ylab = 'acc_train')
par(new = TRUE)
plot(iter,acc_train2,type = 'o',ylim = c(0,1),col = "green",ylab = 'acc_train')
par(new = TRUE)
plot(iter,acc_train3,type = 'o',ylim = c(0,1),col = "blue",ylab = 'acc_train')
par(new = TRUE)
plot(iter,acc_train4,type = 'o',ylim = c(0,1),col = "purple",ylab = 'acc_train')
par(new = TRUE)
plot(iter,acc_train5,type = 'o',ylim = c(0,1),col = "yellow",ylab = 'acc_train')
par(new = TRUE)
plot(iter,acc_train6,type = 'o',ylim = c(0,1),col = "orange",ylab = 'acc_train')
par(new = TRUE)
plot(iter,acc_train7,type = 'o',ylim = c(0,1),col = "springgreen",ylab = 'acc_train')
par(new = TRUE)
plot(iter,acc_train8,type = 'o',ylim = c(0,1),col = "royalblue",ylab = 'acc_train')
par(new = TRUE)
plot(iter,acc_train9,type = 'o',ylim = c(0,1),col = "grey",ylab = 'acc_train')
par(new = TRUE)
plot(iter,acc_train10,type = 'o',ylim = c(0,1),col = "pink",ylab = 'acc_train')

```



By choosing different  $C$ s, the training accuracies and the testing accuracies have huge discrepancies. Generally speaking, from the experiment result, we may say the smaller the  $C$  is, the better the performance will be. With a large  $C$ , the SVM is not even able to have a fair performance on classification or regression.

```
rbf <- function(x1,x2,gamma){
  snorm <- sum((x1-x2)^2)
  result <- exp(-gamma*snorm)
  result
}
# use kernel trick
kernel_SVM <- function(X_train, Y_train, X_test, Y_test, lambda = 0.01, gamma = 0.5, num_iterations = 10000){
  {
    n <- dim(X_train)[1]
    #p <- dim(X_train)[2] + 1
    #X_train1 <- cbind(rep(1,n), X_train)
    K_matrix_train <- matrix(0, n, n)
    for (i in 1:n){
      for (j in 1:n){
        iden <- rbf(X_train[i,],X_train[j,],gamma)
        K_matrix_train[i,j] <- iden
      }
    }

    Y_train <- 2*Y_train - 1#0,1->-1,1
    beta <- matrix(rep(0,n), nrow = n)

    ntest <- nrow(X_test)
    #X_test1 <- cbind(rep(1,ntest), X_test)
    Y_test <- 2*Y_test - 1#0,1->-1,1
    K_matrix_test <- matrix(0, ntest, n)
    for (i in 1:ntest){
      for (j in 1:n){
        iden_test <- rbf(X_test[i,],X_train[j,],gamma)

```

```

    K_matrix_test[i,j] <- iden_test
  }
}


acc_train <- rep(0, num_iterations)
acc_test <- rep(0, num_iterations)

for(it in 1:num_iterations){
  s <- K_matrix_train %*% beta
  db <- s*Y_train < 1
  dbeta <- matrix(rep(1,n), nrow = 1)%*%((matrix(db*Y_train, n, n)*K_matrix_train))/n
  beta <- beta +learning_rate *t(dbeta)
  beta[1:n] <- beta[1:n] - lambda *beta[1:n]


  acc_train[it] <- mean(sign(s * Y_train))
  acc_test[it] <- mean(sign(K_matrix_test %*% beta * Y_test))
}
model <- list(beta = beta, acc_train = acc_train, acc_test = acc_test)
model
}

kernel_model <- kernel_SVM(training_digits, training_labels, testing_digits, testing_labels, num_iterat.
kernel_beta <- kernel_model[[1]]
kernel_acc_train <- kernel_model[[2]]
kernel_acc_test <- kernel_model[[3]]
iter <- 1:200
plot(iter, kernel_acc_train, type = 'o',ylim = c(0,1),col = 'red')

```




```
plot(iter, kernel_acc_test, type = 'o', ylim = c(0,1), col = 'green')
```




```
gamma_model <- kernel_SVM(training_digits, training_labels, testing_digits, testing_labels, num_iterati  
gamma_beta <- gamma_model[[1]]  
gamma_acc_train <- gamma_model[[2]]  
gamma_acc_test <- gamma_model[[3]]  
iter <- 1:200
```

```
plot(iter, gamma_acc_train, type = 'o', ylim = c(0,1), col='red')
```




```
plot(iter , gamma_acc_test, type = 'o',ylim = c(0,1),col='green')
```




By changing different gamma of the kernel, we also get different accuracy. Larger gamma makes the model looser.

```
learning_model <- my_SVM(training_digits, training_labels, testing_digits, testing_labels, num_iteration
learning_beta <- learning_model[[1]]
learning_acc_train <- learning_model[[2]]
learning_acc_test <- learning_model[[3]]
iter <- 1:200
plot(iter , learning_acc_train, type = 'o',ylim = c(0,1),col='red')
```





```
plot(iter , learning_acc_test, type = 'o',ylim = c(0,1), col = 'green')
```



By changing learning rate to 1, we find the accuracy converges slower than  $lr=0.1$ , while the path is also not stable. But finally, the accuracy comes to a fair result. But such a large  $lr$  is not a great choice.