

Lec 13: Support Vector Machine (SVM) II

Ailin Zhang

2022-10-24

Agenda

- Dual form
- Kernel SVM
- Linear Inseparability
- Hinge Loss
- Connection to logistic regression

Support Vector Machine

Let u be a unit vector that has the same direction as β . $u = \frac{\beta}{|\beta|}$.

Suppose X_i is an example on the margin (i.e., support vector), the projection of X_i on u is

$$\langle X_i, u \rangle = \langle X_i, \frac{\beta}{|\beta|} \rangle = \frac{X_i^\top \beta}{|\beta|} = \frac{\pm 1}{|\beta|}.$$

So the margin is $1/|\beta|$. In order to maximize the margin, we should minimize $|\beta|$ or $|\beta|^2$. Hence, the SVM can be formulated as an optimization problem as follows:

$$\begin{aligned} & \text{minimize} && \frac{1}{2}|\beta|^2, \\ & \text{subject to} && y_i X_i^\top \beta \geq 1, \forall i. \end{aligned}$$

Recall $X_i^\top \beta$ is the score, and $y_i X_i^\top \beta$ is the individual margin of observation i . This is the **primal form** of SVM.


Dual Form: Lagrange Multiplier

Let $\alpha = (\alpha_1, \alpha_2, \dots, \alpha_n)$, where $\alpha_i \geq 0$

$$L(\beta, \alpha) = \frac{1}{2}|\beta|^2 + \sum_{i=1}^n \alpha_i(1 - y_i X_i^\top \beta)$$

The idea is to solve an unconstrained problem because it is easier to solve.

Dual Form: Lagrange Multiplier and saddle point



Dual Form

$$L(\beta, \alpha) = \frac{1}{2}|\beta|^2 + \sum_{i=1}^n \alpha_i (1 - y_i X_i^\top \beta)$$

1. $\frac{\partial L}{\partial \beta} = 0$

$$\hat{\beta} = \sum_{i=1}^n \alpha_i y_i X_i$$

2. Dual function:

$$Q(\alpha) = \sum_{i=1}^n \alpha_i - \frac{1}{2} \left| \sum_{i=1}^n \alpha_i y_i X_i \right|^2$$

Dual Problem: $\max_{\alpha_i \geq 0} Q(\alpha)$

Dual Form

$$L(\beta, \alpha) = \frac{1}{2}|\beta|^2 + \sum_{i=1}^n \alpha_i (1 - y_i X_i^\top \beta)$$

1. $\frac{\partial L}{\partial \beta} = 0$

$$\hat{\beta} = \sum_{i=1}^n \alpha_i y_i X_i$$

2. Dual function:

$$Q(\alpha) = \sum_{i=1}^n \alpha_i - \frac{1}{2} \left| \sum_{i=1}^n \alpha_i y_i X_i \right|^2$$

Dual Problem: $\max_{\alpha_i \geq 0} Q(\alpha)$ Solve this by coordinate descent


Coordinate Descent

- Each iteration
 - For i in 1 to n :
 $\max_{\alpha_i} Q(\alpha)$ by fixing the rest $\alpha_j, j \neq i$
Remark: All $\alpha_i \geq 0$
- Until Convergence

For prediction: $\hat{y} = \text{sign}(\langle x, \hat{\beta} \rangle)$

Dual Form

The primal form of SVM is max margin, and the dual form of SVM is min distance.



$$\text{max margin} = \text{min distance}$$

The margin between the two sets is defined by the minimum distance between two.

Dual Form - Convex Hull

Let $X_+ = \sum_{i \in +} c_i X_i$ and $X_- = \sum_{i \in -} c_i X_i$
($c_i \geq 0, \sum_{i \in +} c_i = 1, \sum_{i \in -} c_i = 1$) be two points in the positive and negative convex hulls. The margin is $\min |X_+ - X_-|^2$.

$$\begin{aligned} |X_+ - X_-|^2 &= \left| \sum_{i \in +} c_i X_i - \sum_{i \in -} c_i X_i \right|^2 \\ &= \left| \sum_i y_i c_i X_i \right|^2 \\ &= \sum_{i,j} c_i c_j y_i y_j \langle X_i, X_j \rangle, \end{aligned}$$

$$\text{subject to } c_i \geq 0, \sum_{i \in +} c_i = 1, \sum_{i \in -} c_i = 1.$$

Dual Form - Convex Hull

Let $X_+ = \sum_{i \in +} c_i X_i$ and $X_- = \sum_{i \in -} c_i X_i$
($c_i \geq 0, \sum_{i \in +} c_i = 1, \sum_{i \in -} c_i = 1$) be two points in the positive and negative convex hulls. The margin is $\min |X_+ - X_-|^2$.


$$\begin{aligned} |X_+ - X_-|^2 &= \left| \sum_{i \in +} c_i X_i - \sum_{i \in -} c_i X_i \right|^2 \\ &= \left| \sum_i y_i c_i X_i \right|^2 \\ &= \sum_{i,j} c_i c_j y_i y_j \langle X_i, X_j \rangle, \end{aligned}$$

$$\text{subject to } c_i \geq 0, \sum_{i \in +} c_i = 1, \sum_{i \in -} c_i = 1.$$

We can play the kernel trick to replace $\langle X_i, X_j \rangle$ by $K(X_i, X_j)$ Solvable with sequential minimal optimization

Kernel SVM


For a linearly non-separable dataset:



A popular kernel:

- Gaussian radial basis function $K(X, X') = \exp(-\gamma|X - X'|^2)$

Linear Inseparability



we have a few examples that are incorrectly classified. We'd like to somehow move the bad examples to the other side of the hyperplane. But for this, we'd have to pay a price.

$$\begin{aligned} & \text{minimize} && \frac{1}{2}|\beta|^2 + C \sum_{i=1}^n \xi_i, \\ & \text{subject to} && y_i X_i^\top \beta \geq 1 - \xi_i, \forall i. \end{aligned}$$

Slack Variable

$$\begin{aligned} & \text{minimize} && \frac{1}{2}|\beta|^2 + C \sum_{i=1}^n \xi_i, \\ & \text{subject to} && y_i X_i^\top \beta \geq 1 - \xi_i, \forall i. \end{aligned}$$

Essentially, ξ_i is the amount which we move example i , and C is some positive constant.

Dual form:

$$L(\beta, \xi, \alpha, \mu) = \frac{1}{2}|\beta|^2 + C \sum_{i=1}^n \xi_i + \sum_{i=1}^n \alpha_i (1 - \xi_i - y_i X_i^\top \beta) + \sum_{i=1}^n \mu_i (-\xi_i)$$

$$\max_{\alpha, \mu} \min_{\beta, \xi} L(\beta, \xi, \alpha, \mu)$$

$$\min_{\beta, \xi} L(\beta, \xi, \alpha, \mu)$$

$$\frac{\partial L}{\partial \beta} = 0 \rightarrow \hat{\beta} = \sum_{i=1}^n \alpha_i y_i X_i$$

$$\frac{\partial L}{\partial \xi_i} = 0 \rightarrow \alpha_i = C - \mu_i \leq C$$

$$Q(\alpha) = \sum_{i=1}^n \alpha_i - \frac{1}{2} \left| \sum_{i=1}^n \alpha_i y_i X_i \right|^2$$

Hinge Loss

Another way to interpret ξ_i

- $y_i \beta X_i \geq 1 \rightarrow \xi_i = 0$
- $y_i \beta X_i < 1 \rightarrow \xi_i = 1 - y_i \beta X_i$

$\hat{\xi}_i = \max(0, 1 - y_i \beta X_i)$, this is usually called hinge loss

$$\begin{aligned} \text{minimize} \quad & \frac{1}{2} |\beta|^2 + C \sum_{i=1}^n \xi_i, \\ \rightarrow & \frac{1}{2} |\beta|^2 + C \sum_{i=1}^n \max(0, 1 - y_i \beta X_i) \end{aligned}$$

Recall the loss for perceptron is $\max(0, -y_i X_i^\top \beta)$, which penalizes mistakes or negative margins $y_i X_i^\top \beta$. In comparison, the hinge loss does not only penalize the negative margins $y_i X_i^\top \beta$, it also penalizes margins less than 1.

SVM and ridge logistic regression

Rewrite the

$$\text{loss}(\beta) = \sum_{i=1}^n \max(0, 1 - y_i X_i^\top \beta) + \frac{\lambda}{2} |\beta|^2,$$

we can solve β by gradient descent. The gradient is

$$\text{loss}'(\beta) = - \sum_{i=1}^n 1(y_i X_i^\top \beta < 1) y_i X_i + \lambda \beta,$$

where $1(\cdot)$ is the indicator function.

SVM and ridge logistic regression

Rewrite the

$$\text{loss}(\beta) = \sum_{i=1}^n \max(0, 1 - y_i X_i^\top \beta) + \frac{\lambda}{2} |\beta|^2,$$

we can solve β by gradient descent. The gradient is

$$\text{loss}'(\beta) = - \sum_{i=1}^n 1(y_i X_i^\top \beta < 1) y_i X_i + \lambda \beta,$$

where $1(\cdot)$ is the indicator function.

This is similar to the ridge logistic regression

$$\text{loss}(\beta) = \sum_{i=1}^n \log[1 + \exp(-y_i X_i^\top \beta)] + \frac{\lambda}{2} |\beta|^2,$$

$$\text{loss}'(\beta) = - \sum_{i=1}^n \sigma(-y_i X_i^\top \beta) y_i X_i + \lambda \beta.$$

R code for SVM

```
my_SVM <- function(X_train, Y_train, X_test, Y_test, lambda = 0.01,
                   num_iterations = 1000, learning_rate = 0.1)
{
  n      <- dim(X_train)[1]
  p      <- dim(X_train)[2] + 1
  X_train1 <- cbind(rep(1, n), X_train)
  Y_train <- 2 * Y_train - 1
  beta   <- matrix(rep(0, p), nrow = p)

  ntest  <- nrow(X_test)
  X_test1 <- cbind(rep(1, ntest), X_test)
  Y_test <- 2 * Y_test - 1

  acc_train <- rep(0, num_iterations)
  acc_test  <- rep(0, num_iterations)

  for(it in 1:num_iterations)
  {
    s      <- X_train1 %%% beta
    db     <- s * Y_train < 1
    dbeta  <- matrix(rep(1, n), nrow = 1) %%% ((matrix(db*Y, n, p)*X1))/n;
    beta   <- beta + learning_rate * t(dbeta)
    beta[2:p] <- beta[2:p] - lambda * beta[2:p]

    acc_train[it] <- mean(sign(s * Y_train))
    acc_test[it]  <- mean(sign(X_test1 %%% beta * Y_test))
  }
  model <- list(beta = beta, acc_train = acc_train, acc_test = acc_test)
  model
}
```